
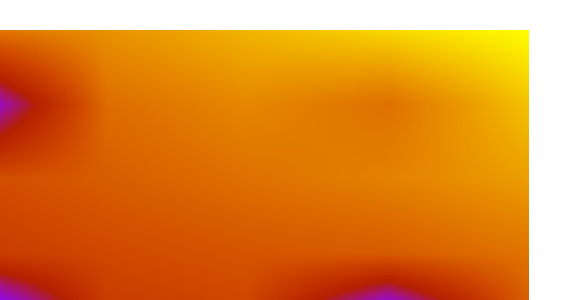


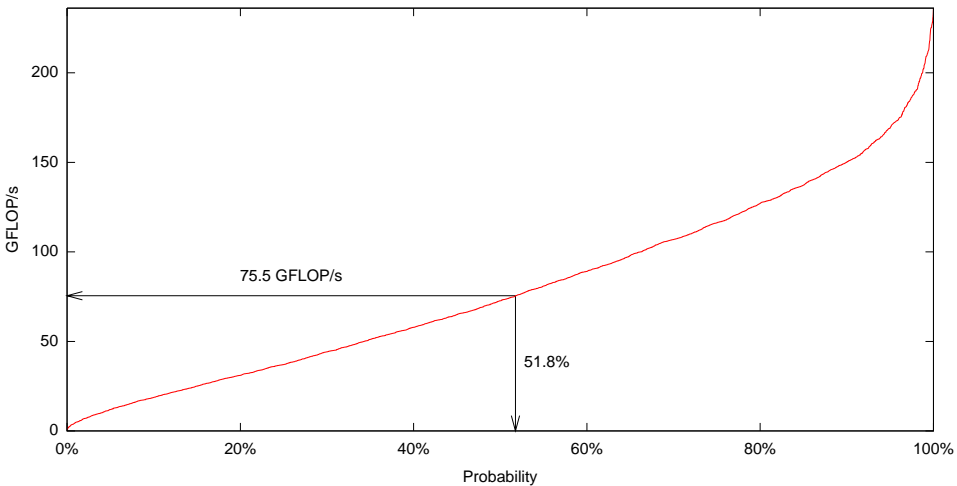
Dgemm/MKL/SMM/Xeon_E5-2670 - Performance





A heatmap showing the number of GFLOPs (Giga Floating Point Operations per second) as a function of the number of layers M (x-axis, ranging from 1 to 9) and the number of heads N (y-axis, ranging from 1 to 9). The color scale ranges from 10 (dark purple) to 80 (yellow). The GFLOPs increase significantly as M increases, especially for $N \geq 6$, where the values reach up to 80 GFLOPs. For $N=1$, the GFLOPs are much lower, around 10-20 GFLOPs.

Dgemm/MKL/SMM/Xeon_E5-2670 - Performance (CDF)



Geo. Mean: 60.7 GFLOP/s Median: 72.7 GFLOP/s

Dgemm/MKL/SMM/Xeon_E5-2670 - Performance (M = N = 4)

